How the *model weights* are split over cores

Data Parallelism

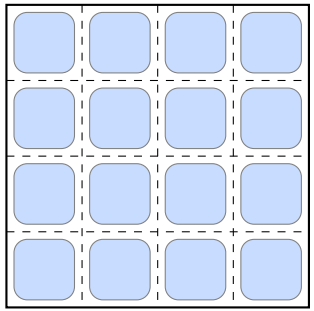

Model Parallelism

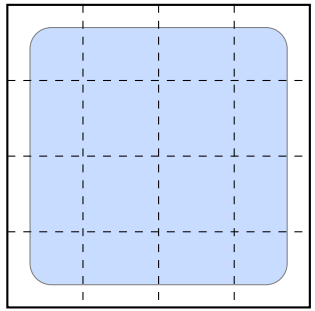

Model and Data Parallelism

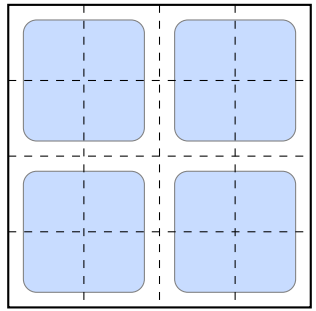

Expert and Data Parallelism

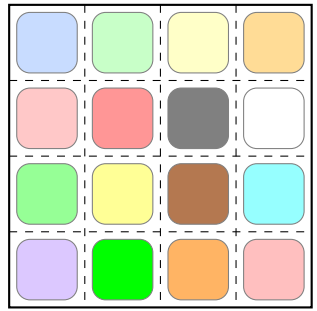

Expert, Model and Data Parallelism

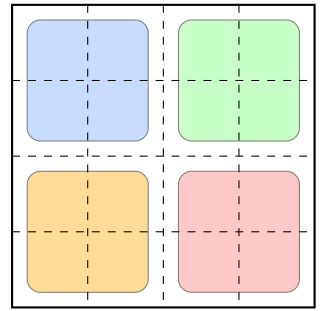

How the *data* is split over cores

Data Parallelism

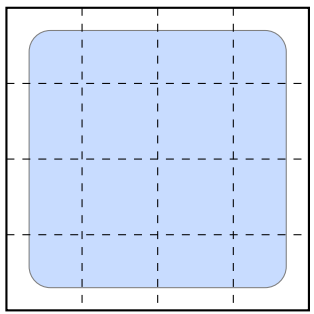

Model Parallelism

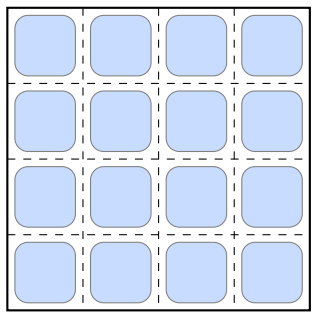

Model and Data Parallelism

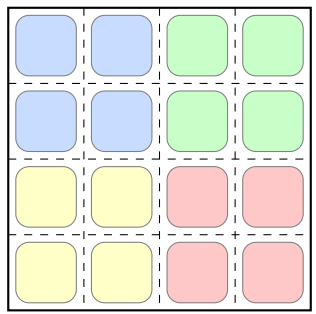

Expert and Data Parallelism

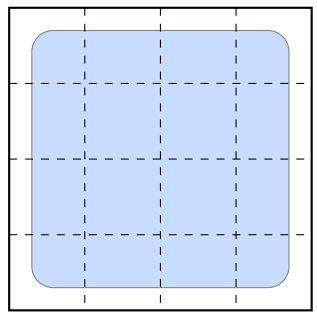

Expert, Model and Data Parallelism